


A cluster of various spheres in white, gold, and blue with gold and white stripes, arranged in a group on the left side of the slide.

Computer Organization

Lab7 Floating-Point Number Processing

A blue pill-shaped button with a small orange circle on its left side.

Floating-Point





Topics

➤ Floating-Point Number

- ✓ IEEE 754 On Single-Precision Floating-Point Number
- ✓ IEEE 754 On Double-Precision Floating-Point Number
- ✓ Conversion between Floating-Point and decimal Number
- ✓ IEEE 754 Single-Precision Floating-Point Number Classification
- ✓ Infinite vs NaN (Floating-Point)

➤ Floating-Point Instructions


- ✓ Floating-Point Registers in RISC-V
- ✓ Floating-Point Instructions Classification
- ✓ Floating-Point system Calls in Rars

➤ Practice



IEEE 754 On Single-Precision Floating-Point Number

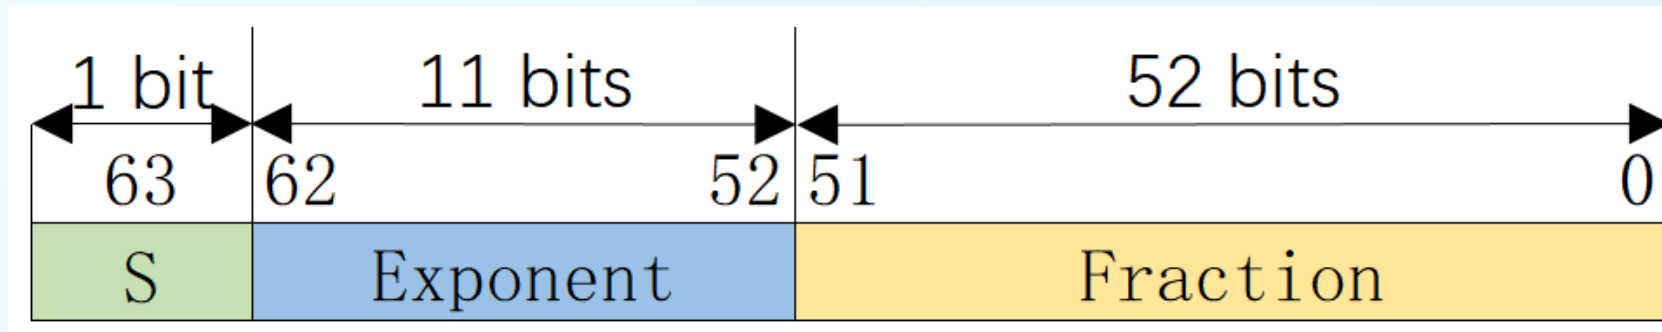
- Single-Precision Floating-Point Number (32-bit width)
 - ✓ Signed bit (符号位): the most significant bit
 - 0: positive number
 - 1: negative number
 - ✓ Exponent (阶码): 8 bits
 - Used to represent indices with a base of 2
 - Using frame shift code (移码), and the **Bias** is 127 (0111_1111)
 - The real value is: **Exponent - Bias**
 - 0000_0000 and 1111_1111 are reserved
 - ✓ Fraction (尾数): 23 bits
 - Representing the decimal part under binary
 - Using true code (原码), and the real value is: **1. Fraction**
 - ✓ The real value of the floating-point number
 - $X = (-1)^S \times (1.\text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$





IEEE 754 On Double-Precision Floating-Point Number

- Double-Precision Floating-Point Number (64-bit width)





Conversion between Floating-Point and decimal Number (1)

➤ Convert 408.6875_{ten} to IEEE 754 single-precision floating-point number

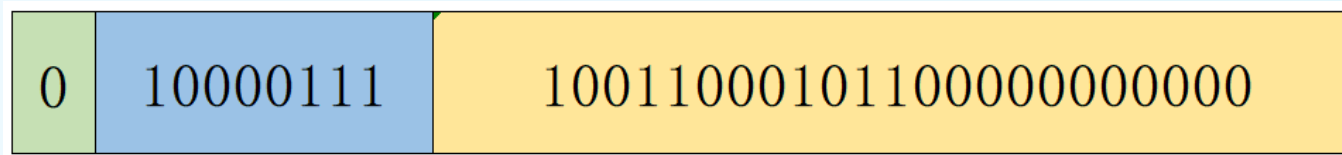
✓ Binary: $110011000.1011 \longrightarrow 2^8 \times 1.100110001011$

✓ Signed bit: 0

✓ Exponent: $8 + 127 = 135_{\text{ten}} = 10000111_{\text{two}}$

✓ Fraction: 100110001011

✓ Floating number:



✓ Convert binary to hexadecimal: 0x43CC5800



Conversion between Floating-Point and decimal Number (2)

- Suppose 0xC1830000 is the hexadecimal machine code for a IEEE 754 single precision floating-point number, convert it to its decimal value.

Binary: 1100 0001 1000 0011 0000 0000 0000 0000

Signed bit: 1

Exponent: $10000011_{\text{two}} = 131_{\text{ten}}$


The real value: $131 - 127 = 4_{\text{ten}}$

Fraction: 000 0011 0000 0000 0000 0000

The real value: 1.0000011

The true value in decimal:

$$\begin{aligned} X &= (-1)^S \times (1.\text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})} \\ &= (-1)^1 \times 1.0000011 \times 2^4 \\ &= -10000.011 \\ &= -16.375 \end{aligned}$$





Conversion between Floating-Point and decimal Number (3)

Piece 7-1

.data

```
fneg1:    .float  -1
wneg1:    .word   -1
fpos1:    .float   1
wpos1:    .word    1
```

Labels	
Label	Address ▲
lab7-piece7-1.asm	
fneg1	0x10010000
wneg1	0x10010004
fpos1	0x10010008
wpos1	0x1001000c

$$\text{➤ } -1 = (-1)^1 \times (1.\text{0}) \times 2^0$$

s: 1; exponent: 0 + 0111_1111; fraction: 0

1	01111111	00000000000000000000000000000000
---	----------	----------------------------------

BF800000

$$\text{➤ } 1 = (-1)^0 \times (1.\text{0}) \times 2^0$$

s: 0; exponent: 0 + 0111_1111; fraction: 0

0	01111111	00000000000000000000000000000000
---	----------	----------------------------------

3F800000

Data Segment				
Address	Value (+0)	Value (+4)	Value (+8)	Value (+c)
0x10010000	0xbf800000	0xffffffff	0x3f800000	0x00000001



IEEE 754 Single-Precision Floating-Point Number Classification

Classification	Sign	Exponent (E)	Fraction (M)	Real value
Positive zero	0	0 (all 0s)	0	+0
Negative zero	1	0 (all 0s)	0	-0
Positive infinity	0	255 (all 1s)	0	$+\infty$
Negative infinity	1	255 (all 1s)	0	$-\infty$
NaN	0 or 1	255 (all 1s)	$M \neq 0$	Not a number
Normalize positive numbers	0	$1 \leq E \leq 254$	M	$(-1)^0 \times (1.M) \times 2^{(E - 127)}$
Normalize negative numbers	1	$1 \leq E \leq 254$	M	$(-1)^1 \times (1.M) \times 2^{(E - 127)}$
Denormalize positive numbers	0	0 (all 0s)	$M \neq 0$	$(-1)^0 \times (0.M) \times 2^{-126}$
Denormalize negative numbers	1	0 (all 0s)	$M \neq 0$	$(-1)^1 \times (0.M) \times 2^{-126}$



Infinite vs NaN (Floating-Point)

- Q1. What are the results of piece 7-2 and 7-3?
- Q2. Which piece of codes will get an infinite value, 7-2 or 7-3?
- Q3. Is the result positive infinity or negative infinity? What's the machine code of this infinity value?
- Q4. Which piece of codes will get the NaN, 7-2 or 7-3? What's the machine code of this NaN value?

```
# Piece 7-2
.include "macro_print_str.asm"
.data
    sdata: .word 0xff7f7fff
    fneg1: .float -1
.text
    la t0, sdata
    flw ft0, (t0)
    fmul.s fa0, ft0, ft0
    li a7, 2
    ecall

    print_string("\n")

    la t0, fneg1
    flw ft0, (t0)
    fmul.s fa0, ft0, ft0
    li a7, 2
    ecall

    li a7, 10
    ecall
```

```
# Piece 7-3
.include "macro_print_str.asm"
.data
    sdata: .word 0xffff7fff
    fneg1: .float -1
.text
    la t0, sdata
    flw ft0, (t0)
    fmul.s fa0, ft0, ft0
    li a7, 2
    ecall

    print_string("\n")

    la t0, fneg1
    flw ft0, (t0)
    fmul.s fa0, ft0, ft0
    li a7, 2
    ecall

    li a7, 10
    ecall
```



Floating-Point Registers in RISC-V

- 32 floating-point registers.
- Each register is 64-bit width.
- View in “Floating Point” window in Rars.

Registers		Floating Point
Name	Number	Value
ft0	0	0x0000000000000000
ft1	1	0x0000000000000000
ft2	2	0x0000000000000000
ft3	3	0x0000000000000000
ft4	4	0x0000000000000000
ft5	5	0x0000000000000000
ft6	6	0x0000000000000000
ft7	7	0x0000000000000000
fs0	8	0x0000000000000000
fs1	9	0x0000000000000000
fa0	10	0x0000000000000000
fa1	11	0x0000000000000000
fa2	12	0x0000000000000000
fa3	13	0x0000000000000000
fa4	14	0x0000000000000000

fa5	15	0x0000000000000000
fa6	16	0x0000000000000000
fa7	17	0x0000000000000000
fs2	18	0x0000000000000000
fs3	19	0x0000000000000000
fs4	20	0x0000000000000000
fs5	21	0x0000000000000000
fs6	22	0x0000000000000000
fs7	23	0x0000000000000000
fs8	24	0x0000000000000000
fs9	25	0x0000000000000000
fs10	26	0x0000000000000000
fs11	27	0x0000000000000000
ft8	28	0x0000000000000000
ft9	29	0x0000000000000000
ft10	30	0x0000000000000000
ft11	31	0x0000000000000000



Floating-Point Instructions Classification

- RV32F & RV32D
- Most instructions start with “f”
- Some basic type
 - ✓ Load and Store: e.g. flw, fld, fsw, fsd
 - ✓ Move data: e.g. fmv.s.x, fmv.x.s
 - ✓ Computational: e.g. fadd.s, fadd.d, fmadd.s, fmadd.d, fmax.d, fmax.s, fsqrt.s, fsqrt.d
 - ✓ Relational: e.g. fle.s, fle.d, flt.s, flt.d
 - ✓ Convert: e.g. fcvt.d.s, fcvt.d.w, fcvt.d.wu
- When the operands are single-precision float number, use “.s” as suffix; when the operands are double-precision float number, use “.d” as suffix.
 - ✓ floating add: fadd.s f1, f2, f3 # assign f1 to f2 + f3
 - ✓ floating add (64 bits): fadd.d f1, f2, f3 # assign f1 to f2 + f3
- Also supports pseudo instructions.



Floating-Point system Calls in Rars

Name	Number	Description	Inputs	Ouputs
PrintFloat	2	Prints a floating point number	fa0 = float to print	N/A
PrintDouble	3	Prints a double precision floating point number	fa0 = double to print	N/A
ReadFloat	6	Reads a float from input console	N/A	fa0 = the float
ReadDouble	7	Reads a double from input console	N/A	fa0 = the double
RandFloat	43	Get a random float	a0 = index of pseudorandom number generator	fa0 = uniformly randomly selected from from [0,1]
RandDouble	44	Get a random double from the range 0.0-1.0	a0 = index of pseudorandom number generator	fa0 = the next pseudorandom



Demo(1)

- Run the codes of piece 7-4, and tell the function of it.
- Change float1 to value of 2147483647.825, what're the outputs, are they correct? Why?

```
# Piece 7-4-1
.include "macro_print_str.asm"
.data
    float1: .float 12.625
    float2: .float 0.5
.text
    la t0, float1
    flw ft0, (t0)
    la t0, float2
    flw ft1, (t0)

    print_string("Original float: ")
    print_float(ft0)

    print_string("\nAfter floor:")
    # floor operation
    fsub.s ft2, ft0, ft1 # ft2 = ft0 - 0.5
    # conver the result to a 32-bit integer
    fcvt.w.s a0, ft2 # a0 = (int32_t)ft2
    li a7, 1
    ecall
```

```
# Piece 7-4-2

    print_string("\nAfter ceil:")
    # ceil operation
    fadd.s ft2, ft0, ft1 # ft2 = ft0 + 0.5
    # conver the result to a 32-bit integer
    fcvt.w.s a0, ft2 # a0 = (int32_t)ft2
    li a7, 1
    ecall

    print_string("\nAfter round:")
    # round operation
    fcvt.w.s a0, ft0 # a0 = (int32_t)ft0
    li a7, 1
    ecall

end
```

```
#Add the content to
"macro_print_str.asm"

.macro print_float(%fr)
    addi sp, sp, -8
    fsw fa0, 4(sp)
    sw a7, 0(sp)

    fmv.s fa0, %fr
    li a7, 2
    ecall

    lw a7, 0(sp)
    flw fa0, 4(sp)
    addi sp, sp, 8
.end_macro
```

```
Original float: 12.625
After floor:12
After ceil:13
After round:13
— program is finished running (0) —
```



Practice 1

- Conversion between hexadecimal floating-point numbers and decimal numbers.
 - ✓ Convert **409.2675_{ten}** to IEEE 754 hexadecimal single precision floating-point numbers.
 - ✓ Suppose **0xc1a6fae1** is the hexadecimal machine code for a IEEE 754 single precision floating-point number, please calculate its corresponding decimal value.
 - ✓ Convert **-409.2675_{ten}** to IEEE 754 hexadecimal double precision floating-point numbers. (optional)
 - ✓ Suppose **0xc0611bf1a9fbe76d** is the hexadecimal machine code for a IEEE 754 double precision floating-point number, please calculate its corresponding decimal value. (optional)
 - ✓ Tip: you can use Rars to get help.



Practice 2-1

- Complete the codes.
- Tip1: There is not an instruction in RISC-V to compare two float numbers and then branch directly. We can use two steps to complete this function.
 - ✓ 1. Compare two float numbers and set an integer register as 1 or 0.
 - ✓ 2. Compare the register and a specific value (0 or 1) to check whether to branch or not.

feq.d t1, f1, f2	Floating Equals (64 bit): if f1 = f2, set t1 to 1, else set t1 to 0
feq.s t1, f1, f2	Floating Equals: if f1 = f2, set t1 to 1, else set t1 to 0
fle.d t1, f1, f2	Floating Less than or Equals (64 bit): if f1 <= f2, set t1 to 1, else set t1 to 0
fle.s t1, f1, f2	Floating Less than or Equals: if f1 <= f2, set t1 to 1, else set t1 to 0
flt.d t1, f1, f2	Floating Less Than (64 bit): if f1 < f2, set t1 to 1, else set t1 to 0
flt.s t1, f1, f2	Floating Less Than: if f1 < f2, set t1 to 1, else set t1 to 0

- Tip 2: When comparing two float numbers, they should both be single-precision or double-precision.

Piece 7-5

```
.include "macro_print_str.asm"
```

```
.data
```

```
fd1: .float 1.0
```

```
dd1: .double 2.0
```

```
.text
```

```
la t0, fd1
```

```
flw ft0, (t0)
```

```
la t0, dd1
```

```
fld ft1, (t0)
```

```
fmv.s fa0, ft0
```

```
li a7, 2
```

```
# print fd1
```

```
ecall
```

```
##complete code here##
```

```
li t2, 1
```

```
beq t1, t2, printLe # if(t1 == 1)
```

```
j printGt
```

```
printLe:
```

```
print_string(" LessOrEqual ")
```

```
j printSecondData
```

```
printGt:
```

```
print_string(" LargerThan ")
```

```
printSecondData:
```

```
fmv.d fa0, ft1
```

```
li a7, 3
```

```
# print dd1
```

```
ecall
```

```
end
```

```
-2.0 LargerThan -80.0
-- program is finished running (0) --

1.0 LessOrEqual 2.0
-- program is finished running (0) --
```



Practice 2-2

- Calculate the value of e from the infinite series:
 - Input a double-precision float number which represents a precision threshold.
 - Your program should terminate when the difference between two successive iterations is smaller than the precision threshold.
 - Print the value of e (as double-precision float).

$$\sum_{n=0}^{\infty} \frac{1}{n!} = \frac{1}{0!} + \frac{1}{1!} + \frac{1}{2!} + \frac{1}{3!} + \frac{1}{4!} + \frac{1}{5!} + \dots$$



Practice 2-3

- Given a single-precision float number 'x' and a positive integer 'r'. Round up 'x' to a number which keeps 'r' digits after the decimal point.

For example, suppose 'x' is 1.5671

if 'r' is 2, print 1.57

if 'r' is 0, print 2

if 'r' is 3, print 1.567