


## Lecture 1

### **Course Introduction**

# 张进



- 2000年-2004年，清华大学电子系学士
- 2004年-2006年，清华大学电子系硕士
- 2006年-2009年，香港科技大学计算机系博士
- 2009年-2014年，香港科技大学计算机系研究助理教授
- 2014年至今，南方科技大学计算机系助理教授、副教授/博士生导师

## 研究方向：移动计算，智能物联网，无线感知，可穿戴设备与移动健康

- 在IOTJ、TWC、TMC、TPDS、Mobicom、Infocom、MobiHoc、ICDCS等国际顶级期刊会议发表论文90余篇
- 近五年承担和参与国家、广东省和深圳市科研项目10余项，包括：国家自然科学基金（青年，面上），科技部重点研发计划，广东省重点实验室，广东省普通高校重点实验室，广东省创新团队，深圳市工程实验室、深圳市学科布局项目
- 带领的研究团队：研究助理教授1名、博士10名、硕士6名、工程师2名、本科生20余名

# Course Introduction

- Today's topics:
  - ◆ Why computer organization is important
  - ◆ Course information
  - ◆ Background and computer abstraction

# Why to Learn Computer Organization?

- Embarrassing if you are a student in CS and can't make sense of the following terms: DRAM, pipelining, cache hierarchies, I/O, virtual memory, ...
- Embarrassing if you are a student in CS and can't decide which processor to buy: 3 GHz P4 or 2.5 GHz Athlon (this course helps us reason about performance/power), ...
- First step for chip designers, compiler/ OS writers
- Knowledge of the hardware will help you write better programs

# Must a Programmer Care about Hardware?

- Must know how to reason about program performance and energy
- CPU Performance: if we understand how CPU process data, we can enhance the computation efficiency
- Memory management: if we understand how/where data is placed, we can help ensure that relevant data is nearby
- Thread management: if we understand how threads interact, we can write smarter multi-threaded programs
- I/O management

# What you have learned?

- Binary numbers
- Read and write basic C/Java programs
- Understand the steps in compiling and executing a program
- Digital Circuit, Logic design:
  - ◆ Logical equations, schematic diagrams
  - ◆ Combinational vs. sequential logic
  - ◆ Finite state machines (FSMs)

# What you will learn?

## ■ Major content

- ◆ Basic parts of a computer (processor, memory, disk, etc.)
- ◆ Principles of computer architecture: CPU datapath and control unit design
- ◆ Assembly language programming in MIPS
- ◆ Memory hierarchies and design
- ◆ I/O organization and design

## ■ Course goals

- ◆ To learn the organizational structures that determine the capabilities and performance of computer systems
- ◆ To understand the interactions between the computer's architecture and its software
- ◆ To understand cost performance trade-offs

# Key Topics

- Introduction (Chapter 1)
  - ◆ Basic terms
  - ◆ Moore's Law, power wall
  - ◆ Core ideas in computer architecture
- Processors (Chapter 2-4)
  - ◆ Assembly language (Chapter 2)
  - ◆ Computer arithmetic (Chapter 3)
  - ◆ Pipelining (Chapter 4)
- Memory (Chapter 5)
- Parallel Processors (Chapter 6)



# The content is useful and important

- Computer organization principles are everywhere
  - ◆ Embedded computer vs. general-purpose computers:
    - Cellphone, Digital Camera, MP3 music player, Industrial process control
- Complex system design
  - ◆ How to partition a problem
  - ◆ Functional Spec → Control & Datapath → Physical implementation
  - ◆ Modern CAD tools
- Both EEs and CSEs need this information in almost all jobs

# Course Information

- Course Materials:

- ◆ Blackboard system: [Computer Organization Spring 2024](#)


- QQ group: 628058404

- Lectures:


- ◆ Friday 14:00-15:50 (3<sup>rd</sup> Teaching Building 305)

- Instructor:

- ◆ Prof. Jin Zhang ([zhangj4@sustech.edu.cn](mailto:zhangj4@sustech.edu.cn))
- ◆ Office: 414, CoE South
- ◆ Office hour: Monday 14:00-16:00



CS202-2024s





# Course Information

## ■ Tentative schedule

WEEK	LECTURE	DATE	TOPIC
1	Lecture #1	Feb. 23, 2024	Introductions
2	Lecture #2	Mar. 1, 2024	RISC-V ISAs: Basics
3	Lecture #3	Mar. 8, 2024	RISC-V ISAs: Procedure Call
4	Lecture #4	Mar. 15, 2024	RISC-V ISAs: Addressing
5	Lecture #5	Mar. 22, 2024	Performance
6	Lecture #6	Mar. 29, 2024	Arithmetic
7	Lecture #7	Apr. 7, 2024	Floating Point Arithmetic
8	Lecture #8	Apr. 12, 2024	The Processor
9	Mid-term	Apr. 19, 2024	TBD (Mid-term exam covers Lecture #1—#8)
10	Lecture #9	Apr. 26, 2024	The Pipeline
11	Break	May 3, 2024	Labors Day Holiday
12	Lecture #10	May 10, 2024	Instruction-Level Parallelism
13	Lecture #11	May. 17, 2024	Memory Hierarchy
14	Lecture #12	May. 24, 2024	Memory Hierarchy(cont.)
15	Lecture #13	May. 31, 2024	Memory Hierarchy(cont.)
16	Lecture #14	Jun. 7, 2024	Parallel Processors

# Course Information

- Textbook:
  - ◆ Computer Organization & Design, the Hardware/Software Interface, **RISC-V edition**. D. A. Patterson and J. L. Hennessy
    - The Textbook uses RV64, in class we learn RV32
- Reference book:
  - ◆ Computer Architecture - a quantitative approach, Hennessy and Patterson, 5<sup>th</sup> edition



# Patterson and Hennessy



- Turing award 2017

For pioneering a systematic, quantitative approach to the design and evaluation of **computer architectures** with enduring impact on the **microprocessor industry**.



David A. Patterson  
Professor of UC Berkeley  
Distinguished Engineer at Google



John L. Hennessy  
President of Stanford University  
Chairman of Alphabet

# Recommended Reading

- Code: The Hidden Language of Computer Hardware and Software, Charles Petzold, 2000. 《编码的奥秘》, 《编码——隐匿在计算机软硬件背后》
- Computer Systems: A Programmer's Perspective, R. E. Bryant, D. R. O'Hallaron, 3<sup>rd</sup> Edition, Pearson, 2015. 《深入理解计算机系统》

# Other Readings

- ❖ 《浪潮之巅》《数学之美》吴军
- ❖ 《创新者：一群技术狂人和鬼才程序员如何改变世界》  
沃尔特·艾萨克森（《史蒂夫·乔布斯传》、《爱因斯坦传》、《本杰明·富兰克林传》）
- ❖ 《沸腾十五年》《沸腾新十年》林军 《激荡四十年》
- ❖ 《图灵和ACM图灵奖》
- ❖ 《黑客与画家》硅谷创业之父Paul Graham
- ❖ 《人月神话》《信息简史》
- ❖ KK三部曲 《失控》《科技想要什么》《必然》

# Course Information

- Assessment:

- ◆ Theoretical part:

- Final exam (~30%)
    - Midterm exam (~30%)
    - Assignments and Attendance (~10%)

- ◆ Lab part (~30%)

- Please Submit on time



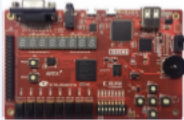
- ◆ all the assignments and reports should be submitted in the BlackBoard system, late submission will not be accepted in the system.



# Labs

- Labs are a key portion of the class
- Highly recommended to find your partner as soon as possible.

## Toolkits used in our Labs

Task	Tool kits
Learn and practice <b>RISC-V</b> ( a type of Assembly language)	➤ <b>Rars (rars_27a7c1f)</b> 
Design and implement an <b>CPU</b>	➤ <b>Vivado</b>  ➤ <b>FPGA based Development Board EGO1</b> 
Test the <b>CPU</b> with <b>program(s)</b> , both of which are based on <b>RISC-V</b>	➤ <b>Assembler (Rars)</b> ➤ <b>Uart Tools</b> ➤ <b>Vivado</b> ➤ <b>FPGA based Development Board(EGO1)</b>

# Rules about Plagiarism

- No Plagiarism is allowed
  - ◆ If plagiarism on homework or project is found for the first time, **the plagiaristic part is graded as 0** and warning is given to the students
  - ◆ If plagiarism is found for the second time, **the course is graded as 0**
  - ◆ For lab/project report, any sentence that is copied from other paper or article **should cite the original source** as the reference, otherwise, the report is considered as plagiarism
- Submit the commitment letter on Blackboard system before homework #1



# Tips for Attending the Lecture

- To get the best use of lecture
  - ◆ interactive
  - ◆ ask whenever you have question, interrupt whenever you want
  - ◆ ask immediately after the class if you are shy
  - ◆ give me suggestions and feedback frequently
- Get the main idea in class, read the details after class

# The Generations of Computers

- First Generation (1940s - 1950s)
  - ◆ **Vacuum Tubes**
- Second Generation (1950s - 1960s)
  - ◆ **Transistors**
- Third Generation (1960s - 1970s)
  - ◆ **Integrated Circuits**
- Fourth/Now Generation (1970s - Present)
  - ◆ **Microprocessors**
  - ◆ **Artificial Intelligence**

# Old School Computers



# New School Computers

From the Small...






# New School Computers

To the very small...



# New School Computers

To the big...





# Classes of Computers

- Personal computers
  - ◆ General purpose, variety of software
  - ◆ Subject to cost/performance tradeoff
- Server computers
  - ◆ Network based
  - ◆ High capacity, performance, reliability
  - ◆ Range from small servers to building sized

# Classes of Computers

- Supercomputers
  - ◆ High-end scientific and engineering calculations
  - ◆ Highest capability but represent a small fraction of the overall computer market
- Embedded computers
  - ◆ Hidden as components of systems
  - ◆ Stringent power/performance/cost constraints

# Evolvment of Computers

**Petaflop**

**Teraflop**

**Gigaflop**

Flop: float point operation

K M G ...?



# The PostPC Era

- Cloud computing
  - ◆ Warehouse Scale Computers (WSC)
  - ◆ Software as a Service (SaaS)
  - ◆ Portion of software run on a Personal Mobile Device and a portion run in the Cloud
  - ◆ Amazon and Google
- Data centers
  - ◆ Millions of computers connected by off-the-shelf networking devices

## Google Data Centers



# The PostPC Era

- Personal Mobile Device (PMD)
  - ◆ Battery operated
  - ◆ Connects to the Internet
  - ◆ Hundreds of dollars
  - ◆ Smart phones, tablets, electronic glasses




# New Computer Architecture for AI era

- AI and big data requires new computer architecture
  - ◆ More Suitable for deep learning
  - ◆ High requirement on parallel
  - ◆ Low energy
- From CPU to GPU, TPU...
- AI chips
  - ◆ SmartPhone
  - ◆ AlphaGo
  - ◆ Autonomous Driving
  - ◆ ChatGPT
- Dozens of companies dive in this area:
  - ◆ Google, NVIDIA, 华为、地平线、寒武纪、燧原
  - ◆ OpenAI

# Components of a Computer


- Same components for all kinds of computer:
  - ◆ Input Device, Output Device, Memory, Processor (Control, Datapath)





# Teardown of MacBook



- 16" LED-backlit IPS Retina display
- Keyboard and Touch Bar
- 2.6 GHz 6-core Intel Core i7
- 16 GB of 2666 MHz DDR4 SDRAM
- 512 GB SSD
- 100 Watt-hour battery
- Speaker and microphone





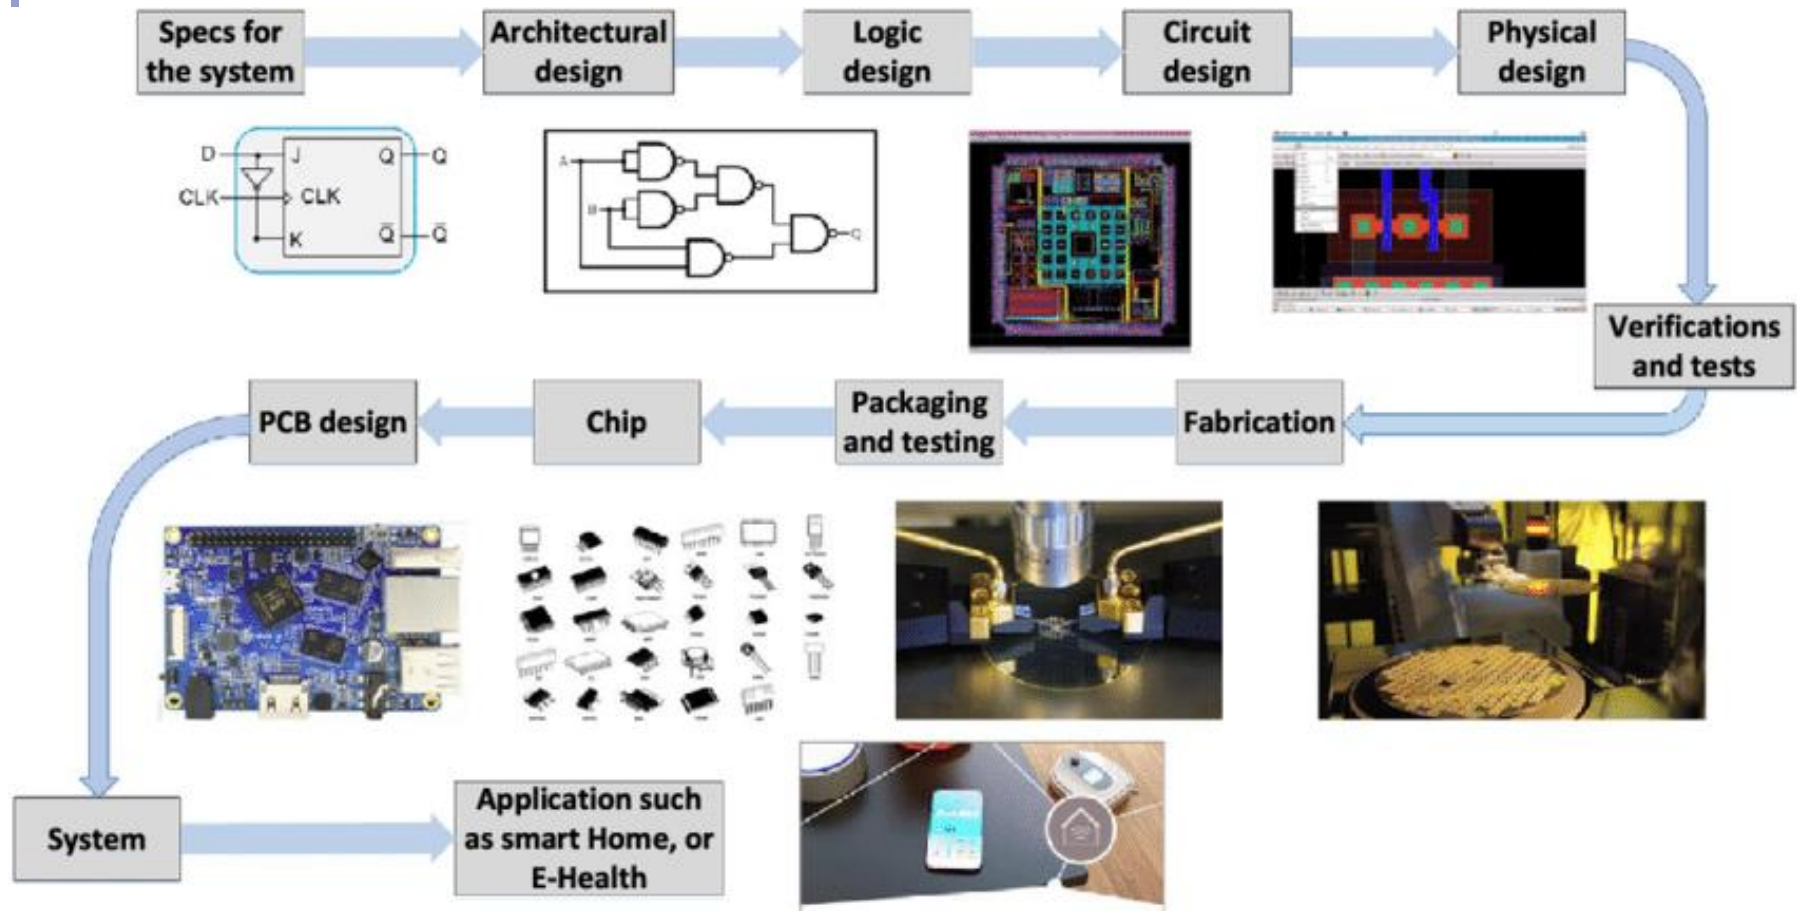
# Inside the Processors

- Datapath:
  - ◆ performs operations on data
- Control:
  - ◆ sequences datapath, memory, I/O
- Cache memory:
  - ◆ small fast Static RAM memory for immediate access to data



Intel's Core i7 wafer and Die map


# Semiconductor manufacturing process chain



# Chip Industry Choke Points

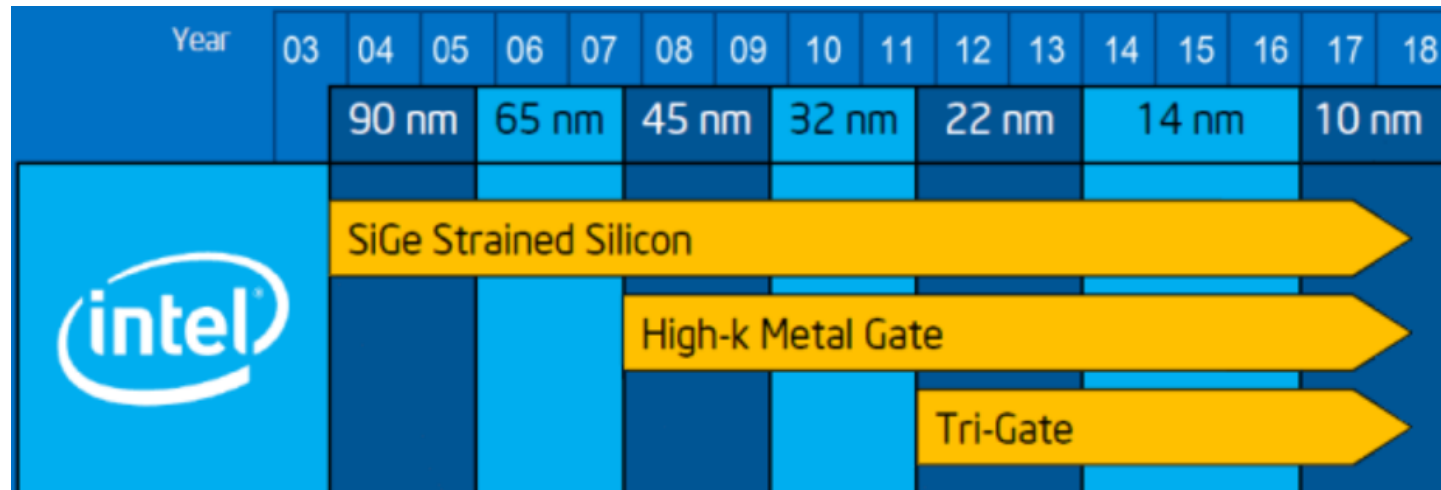
## ■ Key players in chip industry

- ◆ Intel
- ◆ AMD
- ◆ Qualcomm
- ◆ Samsung
- ◆ TSMC
- ◆ Broadcom
- ◆ Nvidia
- ◆ ASML
- ◆ ...



# Processor Technology Trends

- Shrinking of transistor sizes: 90nm(2004) → 45nm(2008) → 22nm(2012) → 10nm(2017) ...

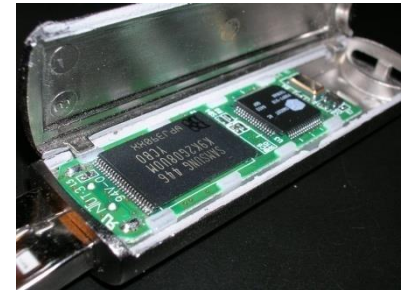


- Transistor density increases by 35% per year and die size increases by 10-20% per year... functionality improvements!
- Transistor speed improves linearly with size (complex equation involving voltages, resistances, capacitances)
- Wire delays do not scale down at the same rate as transistor delays



# Storage

- Volatile main memory
  - ◆ Loses instructions and data when power off
- Non-volatile secondary memory
  - ◆ Magnetic disk
  - ◆ Flash memory
  - ◆ Optical disk (CDROM, DVD)




# Memory and I/O Technology Trends

- DRAM density increases by 40-60% per year, latency has reduced by 33% in 10 years (the memory wall!), bandwidth improves twice as fast as latency decreases
- Disk density improves by 100% every year, latency improvement similar to DRAM
- Networks: primary focus on bandwidth; 10Mb → 100Mb in 10 years; 100Mb → 1Gb in 5 years


# Power Consumption Trends

- Dyn power  $\propto$  activity x capacitance x voltage<sup>2</sup> x frequency
- Voltage and frequency are somewhat constant now, while capacitance per transistor is decreasing and number of transistors (activity) is increasing
- Leakage power is also rising (function of #trans and voltage)



Source: H&P Textbook

# Between Your Program and Hardware



- Application software
  - ◆ Written in high-level language (HLL)
- System software
  - ◆ Compiler: translates HLL code to machine code
  - ◆ Operating System: service code
    - Handling input/output
    - Managing memory and storage
    - Scheduling tasks & sharing resources
- Hardware
  - ◆ Processor, memory, I/O controllers



# Levels of Program Code

- High-level language
  - ◆ Level of abstraction closer to problem domain
  - ◆ Provides for productivity and portability
- Assembly language
  - ◆ Textual representation of instructions
- Hardware representation
  - ◆ Binary digits (bits)
  - ◆ Encoded instructions and data

High-level  
language  
program  
(in C)

```
swap(int v[], int k)
{int temp;
  temp = v[k];
  v[k] = v[k+1];
  v[k+1] = temp;
}
```

Compiler

Assembly  
language  
program  
(for MIPS)

```
swap:
    muli $2, $5, 4
    add  $2, $4, $2
    lw   $15, 0($2)
    lw   $16, 4($2)
    sw   $16, 0($2)
    sw   $15, 4($2)
    jr   $31
```

Assembler

Binary machine  
language  
program  
(for MIPS)



```
000000001010000100000000000011000
000000000000110000001100000100001
100011000110001000000000000000000
100011001111001000000000000000100
101011001111001000000000000000000
101011000110001000000000000000100
00000011111000000000000000001000
```

# Eight Great Ideas

- Design for *Moore's Law*
- Use *abstraction* to simplify design
- Make the *common case fast*
- Performance *via parallelism*
- Performance *via pipelining*
- Performance *via prediction*
- *Hierarchy* of memories
- *Dependability* via redundancy



# Moore's Law



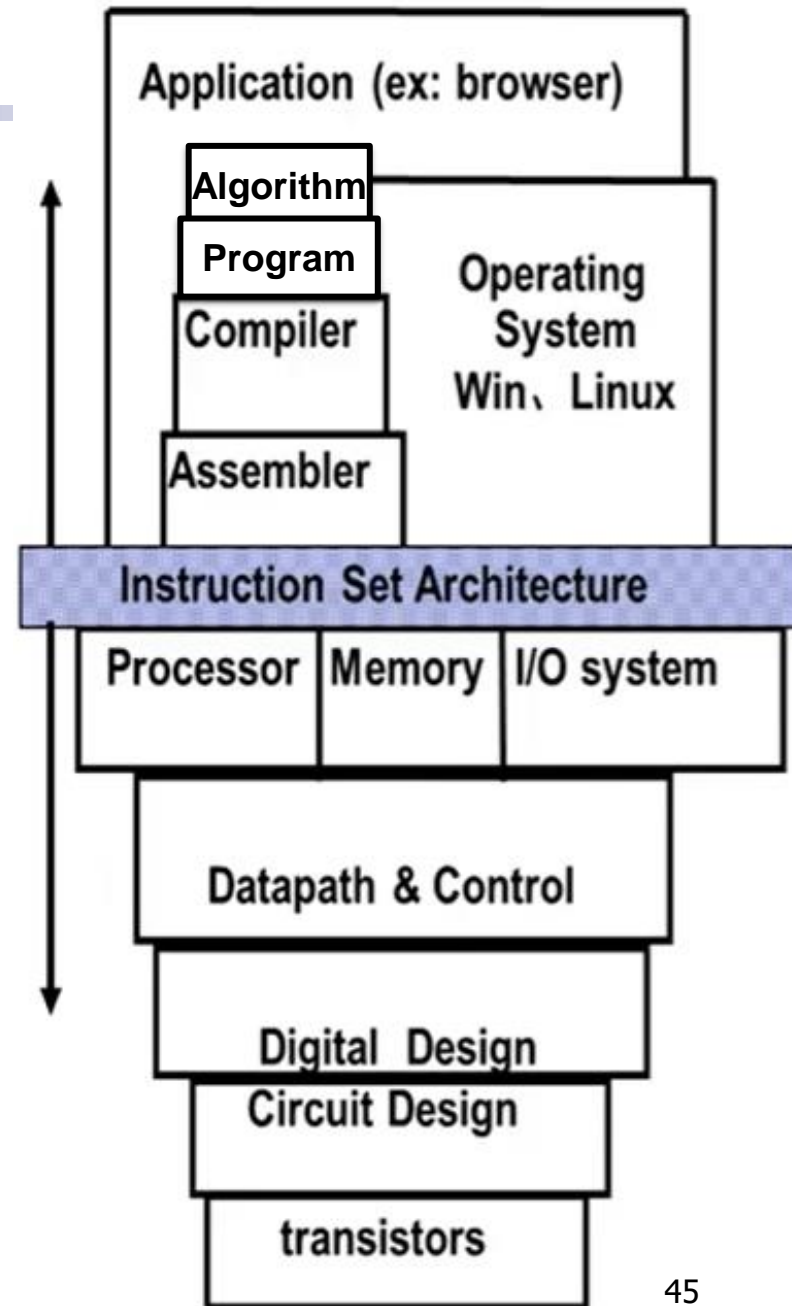
**Gordon Moore  
Intel Cofounder  
B.S. Cal 1950!**

# Post Moore's Law

- Moore's Law meant that the cost of transistors scaled down as technology scaled to smaller and smaller feature sizes.
- And the resulting transistors resulted in increased single-task performance
- But single-task performance improvements hit a brick wall years ago...
- And now the newest, smallest fabrication processes  $< 14\text{nm}$ , might have greater cost/transistor!!!! So, why shrink????

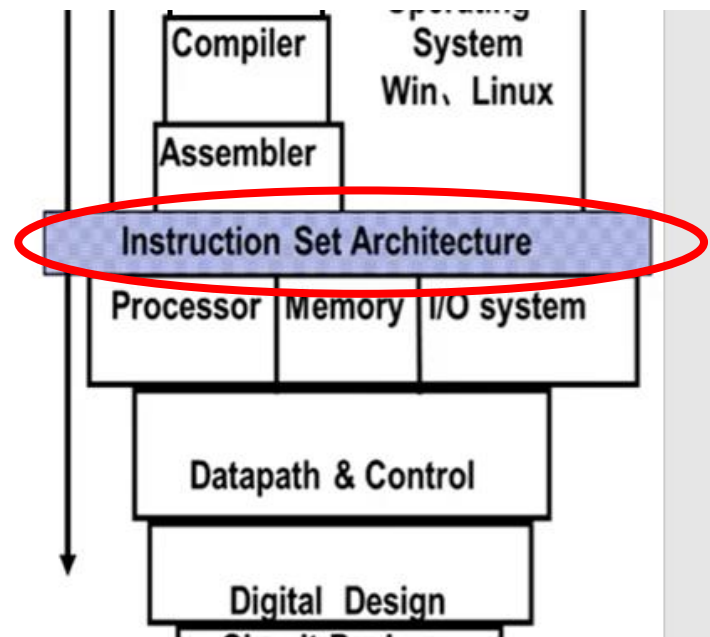
# Abstractions

- Abstraction helps us deal with complexity
  - ◆ Hides lower-level details
- Instruction Set Architecture (ISA) or Computer Architecture
  - ◆ The hardware/software interface
  - ◆ Includes instructions, registers, memory access, I/O, and so on
- Operating system hides details of doing I/O, allocating memory from programmers

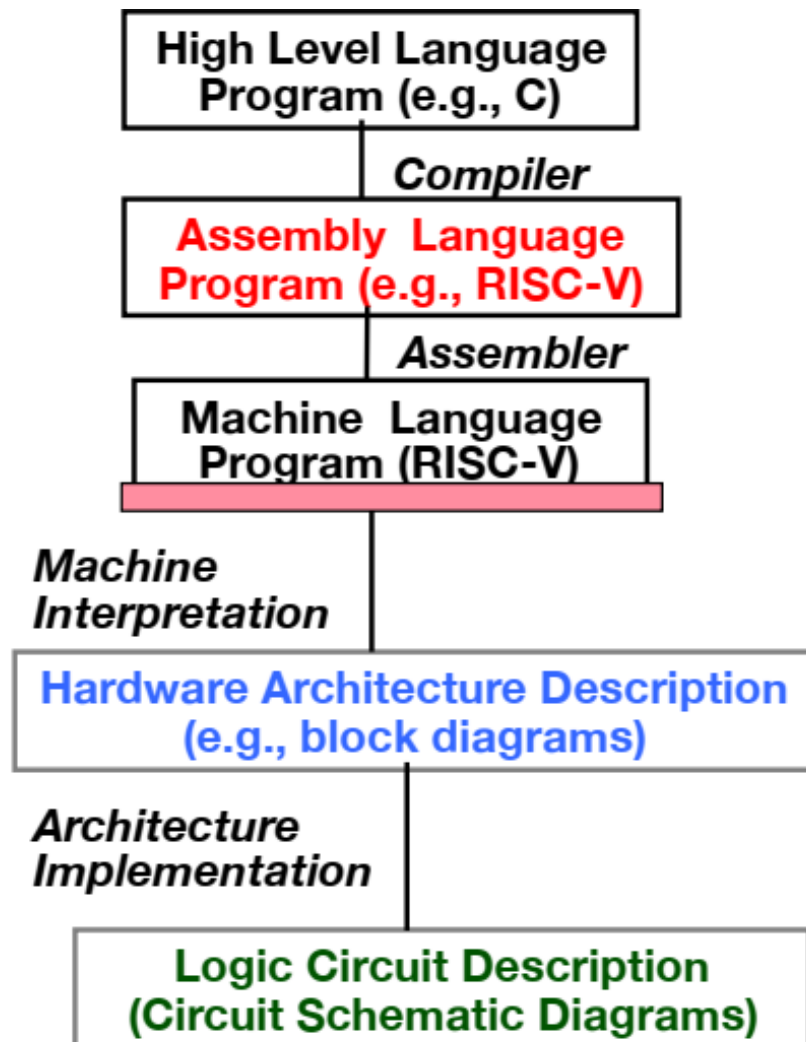


# Instruction Set Architecture (ISA)

- A set of assembly language instructions (ISA) provides a link between software and hardware.
- Given an instruction set, software programmers and hardware engineers work more or less independently.
- Common types of ISA: RISC, CISC
- Examples:
  - ◆ IBM370/X86 (CISC)
  - ◆ **RISC-V** (RISC)
  - ◆ MIPS (RISC)
  - ◆ ARM (RISC)



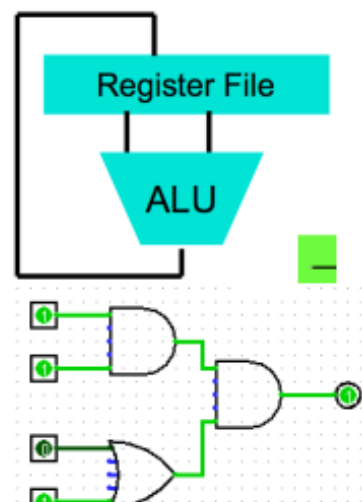
# Abstraction




```
temp = v[k];
v[k] = v[k+1];
v[k+1] = temp;
```

Anything can be represented  
as a *number*,  
i.e., data or instructions

```
0000 1001 1100 0110 1010
1010 1111 0101 1000 0000
1100 0110 1010 1111 0101
0101 1000 0000 1001 1100
```



# Pipeline






# Parallelism

- Raspberry Pi 4 B+
  - Quad core processor at 1.5 GHz
    - Each core is 3-issue, out-of-order superscalar
  - Plus GPU, 1-4 GB RAM, 2x USB3, 2x USB2, Gigabit Ethernet, 2x HDMI...
  - \$35-55
    - Nick is working on a board to turn one of these to power a fully autonomous, vision-guided drone
- Compare with a Cray-1 from 1975:
  - 8 MB RAM, 80 MHz processor, 300MB storage, \$5M+
- Or modern high end servers:
  - You can get a 2u server which supports 4 processors
    - And each processor can have 20+ cores, so 80 processor cores!



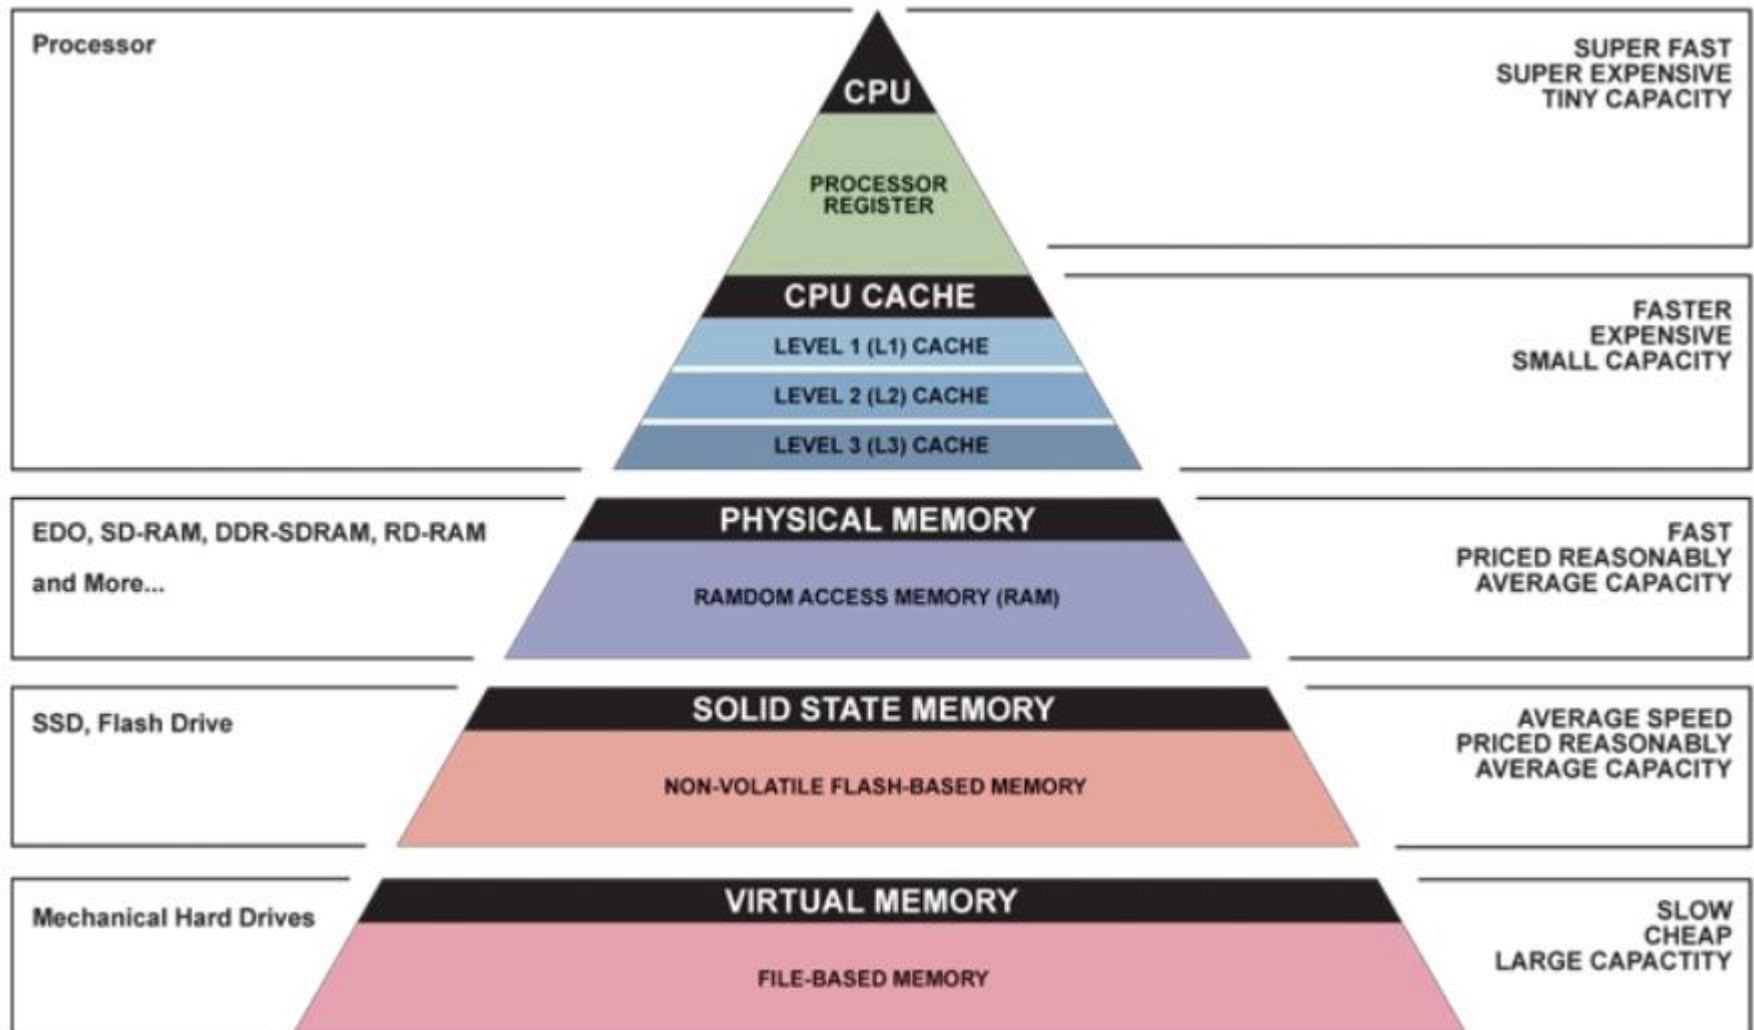
# Jim Gray's Storage Latency Analogy

- How far away is the Data?



**Jim Gray**  
**Turing Award**  
**B.S. Cal 1966**  
**Ph.D. Cal 1969!**

# Memory Hierarchy



# Fails Happen, so?

- 4 disks/server, 50,000 servers
- Failure rate of disks: 2% to 10% / year
  - Assume 4% annual failure rate
- On average, how often does a disk fail?
  - a) 1 / month
  - b) 1 / week
  - c) 1 / day
  - d) 1 / hour


$50,000 \times 4 = 200,000$  disks

$200,000 \times 4\% = 8000$  disks fail

$365 \text{ days} \times 24 \text{ hours} = 8760$  hours

# Reliability via Redundancy

Redundancy so that a failing piece doesn't make the whole system fail



Increasing transistor density reduces the cost of redundancy